

Evaluation Report

Created 2026-08-24T21:56:45



Research-ready metrics

Prompt adherence, image quality, temporal stability and efficiency.

Green = measured | Grey = needs data

Generation metrics

- Generation Time: 45.0 seconds
- Inference Cost: 0.0 USD
- Gpu Usage: Not measured — nvidia-smi is unavailable (common for cloud/API generation).

Image quality

- Resolution: 1024 x 1024 pixels
- Sharpness Variance Laplacian: 621.59 higher is sharper
- Aesthetic Quality Proxy: 89.44 0-100 (Heuristic aesthetic proxy from sharpness, exposure and colour; do not present as a learned aesthetic score.)
- Niqe: 7.0425 lower is better
- Brisque: 60.6245 lower is better
- Psnr: Not measured — A matched --reference-image is required.
- Ssim: Not measured — A matched --reference-image is required.
- Lpips: Not measured — A matched --reference-image is required.

Metric details

continued

Video quality and stability

- Resolution: 960 x 960 pixels
- Duration: 5.031 seconds
- Frame Rate: 32.0 FPS
- Sampled Frames: 16 frames
- Scene Consistency Temporal Ssim: 0.78008 0–1 (Adjacent sampled frames; high values indicate stable appearance, not necessarily realism.)
- Fvd: Not measured — A matched --reference-video is required.
- Temporal Flicker: 0.04451 0-1, lower is better (Mean adjacent-frame luminance change; intended motion can increase this value.)
- Motion Smoothness Optical Flow: 4.76423 pixels/frame, lower is smoother (Mean change in dense optical flow between adjacent samples.)
- Subject Consistency Proxy: 0.68787 0-1, higher is more stable (Central 50% x 60% crop SSIM proxy; assumes the person is centred.)
- Background Consistency Proxy: 0.86757 0-1, higher is more stable (Outer-border SSIM proxy; assumes a mostly fixed camera.)
- Aligned Frame Psnr: Not measured — A matched --reference-video is required.
- Aligned Frame Ssim: Not measured — A matched --reference-video is required.

AI-specific metrics

- Prompt Adherence Clip: 32.791 CLIP logit (Higher means stronger image–prompt alignment; compare within one study.)
- Object Consistency: Not measured — Requires labelled object detections across a test set; do not infer it from a single asset.
- Character Consistency: Not measured — Requires identity annotations or face embeddings across video frames/test cases.
- Hallucination Rate: Not measured — Requires a labelled set of requested objects versus generated objects.

Methodology note

- Interpretation: No-reference scores evaluate the generated asset itself. PSNR, SSIM, LPIPS and FVD need valid matched references or a dataset. Report mean and standard deviation over at least 30 prompts.